

Report_Draft_PrepA_v5__

December 5, 2024

1 Data Collation, Exploration & Preparation

```
[1]: # General setup and imports used throughout the Jupyter Notebook
#
# Libraries For file handling and dataframes
import os
import json
from IPython.display import display
import pandas as pd

# Libraires fo plots
import matplotlib.pyplot as plt
plt.style.use('ggplot')

# Other Libraries
import math
import nltk
nltk.download('vader_lexicon')

# Variables used throughout the notebook
DATA_DIRECTORY = 'JCPenney_Data_Original' # Designated data folder within the
↳current working directory
AUGMENTED_DATA = 'Data_Additional' # Additional data sources

# A simple utility function to obtain and summarise key elements of a provided
↳dataframe
#
def print_file_summary(data_frame):
    # Create a temporary df and ensure no lists remain, so that unique items
    ↳can be identified for uniqueness
    temp_df = data_frame.copy()
    temp_df = temp_df.map(lambda cell: str(cell) if isinstance(cell, list) else
    ↳cell)

    # Calculate some
    summary_of_df = pd.DataFrame({'Count': data_frame.count(),
```

```

        'Missing': data_frame.isnull().sum(), 'Empty': 0,
        'Unique': temp_df.nunique(),
        'Type': data_frame.dtypes,
        'String': 0, 'Int': 0, 'Float': 0, 'List': 0
    })

    summary_of_df['Empty'] = (data_frame == '').sum()
    summary_of_df['String'] = data_frame.map(lambda cell: isinstance(cell, str)).sum()
    summary_of_df['Int'] = data_frame.map(lambda cell: isinstance(cell, int)).sum()
    summary_of_df['Float'] = data_frame.map(lambda cell: isinstance(cell, float)).sum()
    summary_of_df['List'] = data_frame.map(lambda cell: isinstance(cell, list)).sum()

    display(summary_of_df)

def print_full_summary(title, data_frame):
    # Print the summary and head for the given dataframe
    # Used frequently in the notebook so created a function to reduce repetition
    print(title)
    print_file_summary(data_frame)
    print(f'First 3 Rows')
    display(data_frame.head(3))

```

```

[nltk_data] Downloading package vader_lexicon to
[nltk_data] /Users/stuartgow/nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!

```

1.1 Provided Data Sources & Content

The provided data sources for this analysis of JC Penney consists of two JSON files and three CSV files: - JSON: jcpenny_products, jcpenny_reviewers - CSV: products, reviews, users

It was not immediately obvious what the relationships between the two types of data were but the json and CSV files appear to be partial duplicates of each other; also the three CSV files hold slightly less information (eg sales price is missing from the csv files). The CSV files appear to be a first attempt to extract data from the json files (eg the json products file has a JSON field holding multiple user reviews and this has looks to have been extracted to prepare the reviews.csv file).

Given the above, the approach used in this analysis was to go back to the 'original' JSON files and work from these but with a sanity check against the three CSV files to make sure no data was missed or inconsistent.

1.1.1 Load JSON Data Files

It is assumed that the data is a snapshot extract of sales information from JCP databases and the bulk of this has been flattened and used to create the jcpenny_products.json file with the

jcpenny_reviewers.json file providing details of individual customers.

The two tables below show the data items and key counts for each file.

```
[2]: # Load the JSON product file and examine the format and content
# NB: Use pandas json load to directly create a dataframe

# Products file source
file_name = 'jcpenny_products.json'
file_path = os.path.join(os.getcwd(), DATA_DIRECTORY, file_name)
if not os.path.isfile(file_path):
    raise Exception(f'File not found: {file_path}')

# File load into a Pandas dataframe, retained and not amended
source_jcp_products_df = pd.read_json(file_path, lines=True)

# Initial look at the file and data fields
print(f'File Summary for: {file_name}')
print_file_summary(source_jcp_products_df)
print(f'First 3 Rows')
display(source_jcp_products_df.head(3))

# Tidy up
del file_name, file_path
```

File Summary for: jcpenny_products.json

	Count	Missing	Empty	Unique	Type	String	Int	\
uniq_id	7982	0	0	7982	object	7982	0	
sku	7982	0	67	6044	object	7982	0	
name_title	7982	0	0	6002	object	7982	0	
description	7982	0	543	5620	object	7982	0	
list_price	7982	0	2166	1037	object	7982	0	
sale_price	7982	0	18	2063	object	7982	0	
category	7982	0	636	1169	object	7982	0	
category_tree	7982	0	636	1997	object	7982	0	
average_product_rating	7982	0	0	153	float64	0	0	
product_url	7982	0	0	7982	object	7982	0	
product_image_urls	7982	0	157	6519	object	7982	0	
brand	7982	0	0	721	object	7982	0	
total_number_reviews	7982	0	0	22	int64	0	7982	
Reviews	7982	0	0	7982	object	0	0	
Bought With	7982	0	0	7982	object	0	0	
		Float	List					
uniq_id	0	0						
sku	0	0						
name_title	0	0						

description	0	0
list_price	0	0
sale_price	0	0
category	0	0
category_tree	0	0
average_product_rating	7982	0
product_url	0	0
product_image_urls	0	0
brand	0	0
total_number_reviews	0	0
Reviews	0	7982
Bought With	0	7982

First 3 Rows

	uniq_id	sku \
0	b6c0b6bea69c722939585baeac73c13d	pp5006380337
1	93e5272c51d8cce02597e3ce67b7ad0a	pp5006380337
2	013e320f2f2ec0cf5b3ff5418d688528	pp5006380337

	name_title \
0	Alfred Dunner® Essential Pull On Capri Pant
1	Alfred Dunner® Essential Pull On Capri Pant
2	Alfred Dunner® Essential Pull On Capri Pant

	description	list_price	sale_price \
0	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16
1	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16
2	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16

	category	category_tree	average_product_rating \
0	alfred dunner	jcpenny women alfred dunner	2.625
1	alfred dunner	jcpenny women alfred dunner	3.000
2	view all	jcpenny women view all	2.625

	product_url \
0	http://www.jcpenny.com/alfred-dunner-essentia...
1	http://www.jcpenny.com/alfred-dunner-essentia...
2	http://www.jcpenny.com/alfred-dunner-essentia...

	product_image_urls	brand \
0	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner
1	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner
2	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner

	total_number_reviews	Reviews \
0	8	[{'User': 'fsdv4141', 'Review': 'You never hav...
1	8	[{'User': 'tpcu2211', 'Review': 'You never hav...
2	8	[{'User': 'pcfg3234', 'Review': 'You never hav...

```

                                Bought With
0  [898e42fe937a33e8ce5e900ca7a4d924, 8c02c262567...
1  [bc9ab3406dcaa84a123b9da862e6367d, 18eb69e8fc2...
2  [3ce70f519a9cfdd85cdbdec358e5347, b0295c96d2b...

```

```

[3]: # Load the JSON reviewers file and examine the format and content
      # NB: Use pandas json load to directly create a dataframe

      # Reviewers file source
      file_name = 'jcpenney_reviewers.json'
      file_path = os.path.join(os.getcwd(), DATA_DIRECTORY, file_name)
      if not os.path.isfile(file_path):
          raise Exception(f'File not found: {file_path}')

      # File load into a Pandas dataframe, retained and not amended
      source_jcp_reviewers_df = pd.read_json(file_path, lines=True)

      # Initial look at the file and data fields
      print(f'File Summary for: {file_name}')
      print_file_summary(source_jcp_reviewers_df)
      print(f'First 3 Rows')
      display(source_jcp_reviewers_df.head(3))

      # Tidy up
      del file_name, file_path

```

File Summary for: jcpenney_reviewers.json

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
Username	5000	0	0	4999	object	5000	0	0	0
DOB	5000	0	0	52	object	5000	0	0	0
State	5000	0	0	57	object	5000	0	0	0
Reviewed	5000	0	0	4030	object	0	0	0	5000

First 3 Rows

	Username	DOB	State	Reviewed
0	bkpn1412	31.07.1983	Oregon	[cea76118f6a9110a893de2b7654319c0]
1	gqjs4414	27.07.1998	Massachusetts	[fa04fe6c0dd5189f54fe600838da43d3]
2	eehe1434	08.08.1950	Idaho	[]

1.2 Working Data Structure - Validation & Augmentation

1.2.1 Data Structure Summary

The five data sources provided were examined and validated in more detail in order to understand the contents and data structure, and to complete any cleaning required. In addition, areas requiring augmentation were identified and additional data was sourced and combined with the original sources. The steps taken, Python code used and outputs are described in the sub-sections below.

The resulting working data structure is summarised below.

Sales

- Details of all sales activity, 7,982 sales records
- Uniquely identified by the key 'uniq_id', in the Pandas dataframe: sales_df
- Sourced from the provided file: jcpenny_products.json
- A relatively small number of sales prices were missing

Customer Sales Reviews

- Details of all customer reviews, 39,063 in total
- Uniquely identified by combined 'uniq_id' + 'customer_id', in the Pandas dataframe: customer_reviews_df
- Sourced from the provided file: jcpenny_products.json
- Major issues with the quality of the reviews and tend so of limited use in analysis. For example, 15,535 reviews appear to be duplicated accross different customers. This could be a data export issue or even the introduction of fake reviews

Customer Details

- A reference list of 5,001 unique JCP customers who have submitted reviews of purchases
- Uniquely identified by the key 'customer_id', in the Pandas dataframe: customers_df
- Sourced from the provided file: jcpenny_reviewers.json
- Major issues with the quality of date of birth information. Appears to be arificially generated and so of limited use in analysis

Stock Details

- A reference list of 1,154 unique stock items
- Uniquely identified by the key 'sku', in the Pandas dataframe: stock_df
- Sourced from the provided file: jcpenny_products.json
- Derived from the 6,044 unique items in the file jcpenny_products.json
- Some issues with basic data differencies for stock but these have been rationalised

States & Territories

- A reference list of the 57 US states and territories, with population and JCP store numbers per state
- Uniquely identified by the key 'state_ISO', in the Pandas states_df
- Sourced from the file: JCP_Stores_State_Collated.csv
- Orginated from the JCP store locator website and US Census Bureau

1.2.2 States & Territories

A reference list for all US states and territories. Contains 57 items (51 states and 6 territories). This the ISO code for later validation of the provided customers data and with population data and JCP store numbers to assist later geographic analysis.

The data was sourced from: - JCP's store locator, see [website](#) - US Census Bureau, see [website](#)

Data Content After review and vailidation the created dataframe, states_df, has 57 unique items. It consists of:

- territory_flag - Indicates whether a state or a territory
- state_ISO - ISO code of the state, territory
- state_name - Name of the state, territory
- population - Population at 2023
- stores_total - Total number of stores at November 2024

Collation & Validation The additional data file, JCP_Stores_State_Collated.csv, was loaded and validated.

```
[4]: # Establish a reference list of states/territories with additional data to
      ↪augment

# Load the states .csv file, exit if do not exist or are invalid
file_path = os.path.join(os.getcwd(), AUGMENTED_DATA,
      ↪'JCP_Stores_State_Collated.csv')
if not os.path.isfile(file_path):
    raise Exception(f"File not found: {file_path}")
states_df = pd.read_csv(file_path)

# Initial look at the file and data fields
print(f'Summary of States - CSV')
print_file_summary(states_df)
print(f'First 3 rows')
display(states_df.head(3))

# Rename column names & set the index on ISO
states_df = states_df.rename(columns={'State or Territory': 'territory_flag',
      ↪'State_ISO': 'state_ISO',
      ↪'State_Name': 'state_name',
      ↪'Population_2023': 'population',
      ↪'Store_Count': 'stores_total'})
#states_df.set_index(keys='state_ISO', inplace=True)

# Convert population to int
states_df['population'] = states_df['population'].str.replace(',', '').
      ↪astype(int)

# Final look at the file and data fields
print(f'Summary of States - CSV')
print_file_summary(states_df)
print(f'First 3 rows')
display(states_df.head(3))

# Tidy up
del file_path

# Provide a simple unique state lookup of ISO for a given name
```

```
def get_state(state_name):
    matched_state = states_df.loc[states_df['state_name'] == state_name]
    if len(matched_state) == 1:
        return matched_state.iloc[0]
    else:
        return None
```

Summary of States - CSV

	Count	Missing	Empty	Unique	Type	String	Int	Float	\
State or Territory	57	0	0	2	object	57	0	0	
State_ISO	57	0	0	57	object	57	0	0	
State_Name	57	0	0	57	object	57	0	0	
Population_2023	57	0	0	57	object	57	0	0	
Store_Count	57	0	0	26	int64	0	57	0	

	List
State or Territory	0
State_ISO	0
State_Name	0
Population_2023	0
Store_Count	0

First 3 rows

	State or Territory	State_ISO	State_Name	Population_2023	Store_Count
0	State	AL	Alabama	5,108,468	9
1	State	AK	Alaska	733,406	1
2	State	AZ	Arizona	7,431,344	17

Summary of States - CSV

	Count	Missing	Empty	Unique	Type	String	Int	Float	\
territory_flag	57	0	0	2	object	57	0	0	
state_ISO	57	0	0	57	object	57	0	0	
state_name	57	0	0	57	object	57	0	0	
population	57	0	0	57	int64	0	57	0	
stores_total	57	0	0	26	int64	0	57	0	

	List
territory_flag	0
state_ISO	0
state_name	0
population	0
stores_total	0

First 3 rows

	territory_flag	state_ISO	state_name	population	stores_total
0	State	AL	Alabama	5108468	9
1	State	AK	Alaska	733406	1

1.2.3 Customers

Details of customers that have completed a review of a purchase made. With 5,001 unique customer records. All customers have a date of birth, however examination of this showed that only 52 were unique and all appear to be artificially generated. Arguably this field should be dropped as it will not provide any meaningful results. However, it has been retained purely so that it can be used to demonstrate analysis techniques.

Data Content After review and validation the created dataframe, `customers_df`, has 5,001 unique customers. It consists of:

- `customer_id` - A unique alphanumeric id
- `DOB` - Date of birth
- `state_ISO` - ISO code for the state or territory. A cross-reference to the `states_df`

Collation & Validation The provided data file, `jcpenny_reviewers.json`, was examined. As these appears to be detailing customers that have completed a review, the term ‘customer’ was used instead of reviewer. The following actions were taken:

- **Fields Rename:** Columns renamed to be consistent with other dataframes
- **Duplicates:** One `customer_id` was used twice. To preserve information, it was decided to keep the duplicates and assign them a new unique `customer_id`
- **Date of Birth:** Surprisingly for 5,000 customers only 52 birth dates were found. Closer examination revealed that a day, month sequence was incremented across years; with the dates range only being from 26 July to 8 August. All converted to a date field
- **States:** When validating against the states reference file to obtain ISO codes, 187 customers did not match due to the incorrect naming of the US Virgin Islands and US Minor Outlying Islands, so these were corrected. Only the ISO code was retained and the full state name dropped, in preference to it being looked up when required
- `uniq_id_list` - This list was dropped once the details had been cross-checked against the new `customer_reviews` and `sales` dataframes.

???? - 4999 in the CSV file what have they done with the duplicate

```
[5]: # Establish a customer details dataframe

# Create a new dataframe for all customer reviews
customers_df = source_jcp_reviewers_df.copy()

# Rename customer column names and validate content for each
customers_df = customers_df.rename(columns={'Username': 'customer_id',
                                           'State': 'state_name',
                                           'Reviewed': 'uniq_id_list'})

# Print the file and data fields
print(f'Summary for customers')
print_file_summary(customers_df)
```

```

print(f'First 3 rows - Renamed Columns')
display(customers_df.head(3))

# Identify duplicate customers
duplicates_flag = customers_df.duplicated(subset=['customer_id'], keep=False)
duplicated = customers_df[duplicates_flag]
print(f'Duplicated Customers:')
display(duplicated)

# Replace duplicates with new customer_id 'DUPnnnnxxxxxx' to preserve
# Use iter tuples as faster for larger datasets
dup_count = 0
for row in duplicated.iteruples():
    dup_count += 1
    new_id = 'DUP' + str(dup_count).zfill(3) + row.customer_id
    customers_df.at[row.Index, 'customer_id'] = new_id

# Double check no duplicates remain
duplicates_flag = customers_df.duplicated(subset=['customer_id'], keep=False)
duplicated = customers_df[duplicates_flag]
print(f'Double-Check No Remaining Duplicated Customers:')
display(duplicated)

# DOB convert to date format and examine the dates used
customers_df['DOB'] = pd.to_datetime(customers_df['DOB'], dayfirst=True,
    ↪errors='coerce')
dates = customers_df.groupby('DOB').size().reset_index(name='counts')
print(f'Dates Count:')
display(dates)
# Drop the date as looks artificially generated and so of no real use in later
    ↪analysis
# customers_df = customers_df.drop('DOB', axis=1)

# States validation - lookup ISO codes, add to customer data and check for
    ↪invalid matches
customers_df['state_ISO'] = customers_df['state_name'].apply(lambda x:
    ↪get_state(x)['state_ISO'] if get_state(x) is not None else None)
unmatched_states = customers_df[customers_df['state_ISO'].isnull()]
print(f'Unmatched States:')
display(unmatched_states[['customer_id', 'state_name']])

# Names mismatch for US Virgin Islands and US Minor Outlying Islands
customers_df.replace('U.S. Virgin Islands', 'US Virgin Islands', inplace=True)
customers_df.replace('Minor Outlying Islands', 'US Minor Outlying Islands',
    ↪inplace=True)

# Repeat the checks & drop state_name if all ISO populated

```

```

customers_df['state_ISO'] = customers_df['state_name'].apply(lambda x:
    ↪get_state(x)['state_ISO'] if get_state(x) is not None else None)
unmatched_states = customers_df[customers_df['state_ISO'].isnull()]
print(f'Unmatched States:')
display(unmatched_states[['customer_id', 'state_name']])

# Drop the state name, rely on the ISO code and states lookup
if len(unmatched_states) != 0:
    raise Exception(f'Cannot match: {len(unmatched_states)} states')
customers_df = customers_df.drop('state_name', axis=1)

# Visual check on state details
states = customers_df.groupby('state_ISO').size().reset_index(name='counts')
print(f'Customers by State:')
display(states)

# Reviewed validate
# TODO: x-check these to sales activity and to reviews to make sure consistent

# Tidy up
del duplicates_flag, duplicated, dup_count, new_id, row
del dates
del unmatched_states, states

```

Summary for customers

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
customer_id	5000	0	0	4999	object	5000	0	0	0
DOB	5000	0	0	52	object	5000	0	0	0
state_name	5000	0	0	57	object	5000	0	0	0
uniq_id_list	5000	0	0	4030	object	0	0	0	5000

First 3 rows - Renamed Columns

	customer_id	DOB	state_name	uniq_id_list
0	bkpn1412	31.07.1983	Oregon	[cea76118f6a9110a893de2b7654319c0]
1	gqjs4414	27.07.1998	Massachusetts	[fa04fe6c0dd5189f54fe600838da43d3]
2	eehe1434	08.08.1950	Idaho	[]

Duplicated Customers:

	customer_id	DOB	state_name	uniq_id_list
731	dqft3311	28.07.1995	Tennessee	[5f280fb338485cfc30678998a42f0a55]
2619	dqft3311	03.08.1969	New Mexico	[571b86d307f94e9e8d7919b551c6bb52]

Double-Check No Remaining Duplicated Customers:

Empty DataFrame

Columns: [customer_id, DOB, state_name, uniq_id_list]

Index: []

Dates Count:

	DOB	counts
0	1950-08-08	99
1	1951-08-08	95
2	1952-08-07	103
3	1953-08-07	112
4	1954-08-07	79
5	1955-08-07	93
6	1956-08-06	96
7	1957-08-06	93
8	1958-08-06	96
9	1959-08-06	94
10	1960-08-05	107
11	1961-08-05	101
12	1962-08-05	106
13	1963-08-05	106
14	1964-08-04	107
15	1965-08-04	106
16	1966-08-04	94
17	1967-08-04	90
18	1968-08-03	91
19	1969-08-03	99
20	1970-08-03	101
21	1971-08-03	90
22	1972-08-02	91
23	1973-08-02	102
24	1974-08-02	102
25	1975-08-02	106
26	1976-08-01	87
27	1977-08-01	97
28	1978-08-01	79
29	1979-08-01	106
30	1980-07-31	99
31	1981-07-31	85
32	1982-07-31	98
33	1983-07-31	99
34	1984-07-30	80
35	1985-07-30	100
36	1986-07-30	83
37	1987-07-30	99
38	1988-07-29	100
39	1989-07-29	81
40	1990-07-29	103
41	1991-07-29	104
42	1992-07-28	101
43	1993-07-28	96
44	1994-07-28	86
45	1995-07-28	95
46	1996-07-27	81

47	1997-07-27	97
48	1998-07-27	111
49	1999-07-27	104
50	2000-07-26	90
51	2001-07-26	80

Unmatched States:

	customer_id	state_name
29	wjfh4432	Minor Outlying Islands
104	ulkz1412	Minor Outlying Islands
106	bsqg4331	Minor Outlying Islands
203	bbiv3413	Minor Outlying Islands
215	surt1311	U.S. Virgin Islands
...
4872	ypcn2342	U.S. Virgin Islands
4940	lric2324	U.S. Virgin Islands
4960	okun1224	Minor Outlying Islands
4970	kjgm1311	U.S. Virgin Islands
4976	gjed1211	U.S. Virgin Islands

[187 rows x 2 columns]

Unmatched States:

Empty DataFrame

Columns: [customer_id, state_name]

Index: []

Customers by State:

	state_ISO	counts
0	AK	94
1	AL	95
2	AR	92
3	AS	86
4	AZ	71
5	CA	99
6	CO	85
7	CT	82
8	DC	83
9	DE	106
10	FL	89
11	GA	79
12	GU	73
13	HI	88
14	IA	94
15	ID	79
16	IL	69
17	IN	86
18	KS	90

19	KY	99
20	LA	80
21	MA	107
22	MD	77
23	ME	94
24	MI	76
25	MN	77
26	MO	84
27	MP	102
28	MS	94
29	MT	97
30	NC	68
31	ND	85
32	NE	90
33	NH	83
34	NJ	101
35	NM	96
36	NV	90
37	NY	83
38	OH	81
39	OK	100
40	OR	96
41	PA	86
42	PR	83
43	RI	93
44	SC	77
45	SD	79
46	TN	89
47	TX	83
48	UM	92
49	UT	80
50	VA	96
51	VI	95
52	VT	103
53	WA	94
54	WI	84
55	WV	80
56	WY	86

1.2.4 Stock Details

Details of all stock (product) data. Contains 1,154 unique lines of stock with each uniquely identified by the key 'sku'. For each stock line the details include description and its list price. There was significant inconsistency of basic details (name, description, list price, image url) for stock items in the provided products file. An attempt has been made to rationalise the data by retaining the most commonly used data items.

Data Content After review and validation the created dataframe, stock_df, has 1,154 unique stock records. It consists of:

- sku - The unique identifier for the stock item
- stock_name - Short name for the stock item
- description - A long description of the stock
- list_price - The standard price for the stock
- stock_image_url - URL for the website image for the stock
- brand - The manufacturer's name for the stock item, eg 'Alfred Dunner'

Collation & Validation The provided data file, jcpenny_products.json, was examined and all stock specific data extracted into the stock_df dataframe. The following actions were taken:

- Missing SKU ids: 67 were missing, so generated ids were added according in line with the most common format structure, to pp600nnnnnnnn
- Drop Fields: Drop all fields that are sales specific: 'uniq_id', 'sale_price', 'category', 'category_tree', 'average_product_rating', 'product_url', 'total_number_reviews', 'Reviews', 'Bought With'
- Stock Name: A significant number of names differed for the same sku. The first name has been retained
- Description: A significant number of descriptions differed for the same sku or were missing. The first name has been retained. However, still 50 had no description so 'No Description Available' was added
- List Price: A significant number of items had different prices for the same stock. So the most common price for each item was used. Even so, 182 stock items do not have a list price
- Stock Image URL: For 170 stock items, the urls did not all match and so the most frequent one was retained *use request to test a sample of URLs*
- Brand: No missing or duplicated, so just copy one
- Rationalise Stock: Retain only a single unique sku record

1.2.5 Reasons for Splitting Sales & Stock Data

The provided file, jcpenny_products.json, appears to contain core stock information and sales specific information. For example a stock list price and a different sales price that varied depending on different sales categories. So stock data has been split out and cleaned in order to be able to more easily analyse stock vs sales data.

It has been assumed that the field sku is the 'Stock Keeping Unit' see [Wikipedia](#) and should be a unique identifier. Therefore all instances of sku have been reviewed and collapsed down into a stock list dataframe, separate from sales activity. The 6,044 product rows have produced 1,154 stock lines.

```
[6]: # Establish a reference list of all product / stock details
      # And also the initial draft of the sales dataframe for further preparation

      # Create an initial new dataframe for all stock details
      stock_df = source_jcp_products_df.copy()

      # Print the file and data fields
```

```

print_full_summary('Summary For Stock/Sales/Product Data - Initial Look',
    ↪stock_df)

# Flag all missing fields for easier checking and replacement
missing_flag = 'Missing'
stock_df.replace('', missing_flag, inplace=True)
stock_df.fillna(missing_flag, inplace=True)

# Count missing and check formats of SKU
sku_formats = {'pp500nnnnnn': r'^pp500\d{6}',
               '1xxxxxx': r'^1\w{6}',
               'enxxxxnnnnnnnn': r'^en\D\d{10}',
               missing_flag: r'Missing'}

counts = {}
filtered = stock_df.copy()
for sku_format, regex_pattern in sku_formats.items():
    matched = stock_df[stock_df['sku'].str.contains(regex_pattern, na=False)]
    counts[sku_format] = len(matched)
    filtered = filtered[~filtered['sku'].isin(matched['sku'])]
print(f'Counts for SKU missing and format types + formats not matching')
display(counts)
display(filtered)

# Generate ids for missing SKU
# Use itertuples as faster for larger datasets
sku_count = 6000000000
missing_sku = stock_df[stock_df['sku'] == missing_flag]
for row in missing_sku.itertuples():
    sku_count += 1
    new_id = 'pp' + str(sku_count)
    stock_df.at[row.Index, 'sku'] = new_id
# Double-check all updated
missing_sku = stock_df[stock_df['sku'] == missing_flag]
display(missing_sku)

# Create an initial new dataframe for all sales details ready for later
    ↪manipulation
sales_df = stock_df.copy()

# Drop non stock columns
columns_not_required = ['uniq_id', 'sale_price', 'category', 'category_tree',
    ↪'average_product_rating',
                        'product_url',
                        'total_number_reviews', 'Reviews', 'Bought With']
stock_df.drop(columns=columns_not_required, inplace=True, errors='ignore')

# Rename the retained columns

```



```

# (nb all listed for documentation purposes)
stock_df = stock_df.rename(columns={'name_title': 'stock_name', 'description': 'description',
                                   'list_price': 'list_price',
                                   'product_image_urls': 'stock_image_url',
                                   'brand': 'brand'})

# Remove duplicated sku ids rows
# Checking consistency of the other columns
sku_duplicated = stock_df.groupby('sku').filter(lambda sku: len(sku) > 1)
print(f'duplicated skus: {len(sku_duplicated)}')
sku_groups_dup = sku_duplicated.groupby('sku')

# Iterate through all grouped sku ids and validate, select the individual
# column values
new_stock_df = stock_df.head(0).copy()
for sku, group in sku_groups_dup:
    new_sku = stock_df.head(0).copy()
    new_sku.at[0, 'sku'] = sku
    # stock name - Just retain the first record
    new_sku.at[0, 'stock_name'] = group['stock_name'].iloc[0]
    # description - Keep the first non-blank
    non_empty = group[group['description'] != missing_flag]
    if(len(non_empty) != 0):
        new_sku.at[0, 'description'] = non_empty['description'].iloc[0]
    else:
        new_sku.at[0, 'description'] = 'No Description Available'
    # list_price
    non_empty = group[group['list_price'] != missing_flag]
    if(len(non_empty) != 0):
        most_frequent = non_empty['list_price'].value_counts().idxmax()
        new_sku.at[0, 'list_price'] = most_frequent
    else:
        new_sku.at[0, 'list_price'] = 0
    # stock_image_url
    non_empty = group[group['stock_image_url'] != missing_flag]
    if(len(non_empty) != 0):
        most_frequent = non_empty['stock_image_url'].value_counts().idxmax()
        new_sku.at[0, 'stock_image_url'] = most_frequent
    else:
        new_sku.at[0, 'stock_image_url'] = 'No URL Available'
    # brand
    non_empty = group[group['brand'] != missing_flag]
    if(len(non_empty) != 0):
        new_sku.at[0, 'brand'] = non_empty['brand'].iloc[0]
    else:

```

```

new_sku.at[0, 'brand'] = 'No Details Available'

# Add the single row to the temporary duplicated stock dataframe
new_stock_df = pd.concat([new_stock_df, new_sku], ignore_index=True)

# Finally, copy the reduced stock list and set all fields to appropriate types
stock_df = new_stock_df.copy()
stock_df['list_price'] = pd.to_numeric(stock_df['list_price'], errors='coerce')

# Print the file and data fields
print_full_summary('Summary For Stock Data - After Cleaning', stock_df)

# Tidy up
del counts, filtered, matched, regex_pattern, sku_format, sku_formats, \
    ↪missing_flag, non_empty
del sku_count, missing_sku, row, new_id, most_frequent
del new_sku, new_stock_df, columns_not_required, sku, sku_duplicated, \
    ↪sku_groups_dup, group

```

Summary For Stock/Sales/Product Data - Initial Look

	Count	Missing	Empty	Unique	Type	String	Int	\
uniq_id	7982	0	0	7982	object	7982	0	
sku	7982	0	67	6044	object	7982	0	
name_title	7982	0	0	6002	object	7982	0	
description	7982	0	543	5620	object	7982	0	
list_price	7982	0	2166	1037	object	7982	0	
sale_price	7982	0	18	2063	object	7982	0	
category	7982	0	636	1169	object	7982	0	
category_tree	7982	0	636	1997	object	7982	0	
average_product_rating	7982	0	0	153	float64	0	0	
product_url	7982	0	0	7982	object	7982	0	
product_image_urls	7982	0	157	6519	object	7982	0	
brand	7982	0	0	721	object	7982	0	
total_number_reviews	7982	0	0	22	int64	0	7982	
Reviews	7982	0	0	7982	object	0	0	
Bought With	7982	0	0	7982	object	0	0	

	Float	List
uniq_id	0	0
sku	0	0
name_title	0	0
description	0	0
list_price	0	0
sale_price	0	0
category	0	0
category_tree	0	0
average_product_rating	7982	0

product_url	0	0
product_image_urls	0	0
brand	0	0
total_number_reviews	0	0
Reviews	0	7982
Bought With	0	7982

First 3 Rows

	uniq_id	sku \
0	b6c0b6bea69c722939585baeac73c13d	pp5006380337
1	93e5272c51d8cce02597e3ce67b7ad0a	pp5006380337
2	013e320f2f2ec0cf5b3ff5418d688528	pp5006380337

	name_title \
0	Alfred Dunner® Essential Pull On Capri Pant
1	Alfred Dunner® Essential Pull On Capri Pant
2	Alfred Dunner® Essential Pull On Capri Pant

	description	list_price	sale_price \
0	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16
1	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16
2	You'll return to our Alfred Dunner pull-on cap...	41.09	24.16

	category	category_tree	average_product_rating \
0	alfred dunner	jcpenny women alfred dunner	2.625
1	alfred dunner	jcpenny women alfred dunner	3.000
2	view all	jcpenny women view all	2.625

	product_url \
0	http://www.jcpenny.com/alfred-dunner-essentia...
1	http://www.jcpenny.com/alfred-dunner-essentia...
2	http://www.jcpenny.com/alfred-dunner-essentia...

	product_image_urls	brand \
0	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner
1	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner
2	http://s7d9.scene7.com/is/image/JCPenny/DP122...	Alfred Dunner

	total_number_reviews	Reviews \
0	8	[{'User': 'fsdv4141', 'Review': 'You never hav...
1	8	[{'User': 'tpcu2211', 'Review': 'You never hav...
2	8	[{'User': 'pcfg3234', 'Review': 'You never hav...

	Bought With
0	[898e42fe937a33e8ce5e900ca7a4d924, 8c02c262567...
1	[bc9ab3406dcaa84a123b9da862e6367d, 18eb69e8fc2...
2	[3ce70f519a9cfdd85cdbdec358e5347, b0295c96d2b...

Counts for SKU missing and format types + formats not matching

```
{'pp500nnnnnn': 7505, '1xxxxxx': 394, 'enxxxxnnnnnnnn': 13, 'Missing': 67}
```

	uniq_id	sku	\
2269	9fa199671d88a2a3cddd06a0dac02763	0903a80	
3984	4875e80ad4e5d0d8970850046a4c8b8c	PP100000902	
7884	6dcebbf40f3195554080edced28d401b	0903a80	

	name_title	\
2269	KitchenAid® Artisan® 5-qt. Stand Mixer KSM150PS	
3984	Alyx® Gauze Print Tank Top or Millennium Pants	
7884	KitchenAid® Artisan® 5-qt. Stand Mixer KSM150PS	

	description	list_price	sale_price	\
2269	The mixer you've always dreamed of. Unique mix...	604.31	423.01	
3984	Missing	Missing	24.1633	
7884	The mixer you've always dreamed of. Unique mix...	Missing	604.31	

	category	category_tree	\
2269	small appliances	jcpenny for-the-home small appliances	
3984	outfits you'll love	jcpenny women outfits you'll love	
7884	wedding	jcpenny wedding	

	average_product_rating	\
2269	2.750	
3984	3.000	
7884	3.125	

	product_url	\
2269	http://www.jcpenny.com/kitchenaid-artisan-5-q...	
3984	http://www.jcpenny.com/alyx-gauze-print-tank-...	
7884	http://www.jcpenny.com/kitchenaid-artisan-5-q...	

	product_image_urls	brand	\
2269	http://s7d9.scene7.com/is/image/JCPenny/09006...	Kitchen Aid	
3984	http://s7d9.scene7.com/is/image/JCPenny/DP032...	Alyx	
7884	http://s7d2.scene7.com/is/image/JCPenny/DP021...	Kitchen Aid	

	total_number_reviews	Reviews	\
2269	8	[{'User': 'vlfw2311', 'Review': 'I dont know w...	
3984	8	[{'User': 'tlim1231', 'Review': 'I was worried...	
7884	8	[{'User': 'lzci4334', 'Review': 'I dont know w...	

	Bought With
2269	[0f09d5de035bbb347c17f55222d9efa4, dae30fb78a6...
3984	[53cf4a9eb003e2b5e9c63722d1011951, 5b7416f4e6a...
7884	[eb8e7f2068b80379afbae5135b280c7b, 44725052ce6...

Empty DataFrame

Columns: [uniq_id, sku, name_title, description, list_price, sale_price, category, category_tree, average_product_rating, product_url, product_image_urls, brand, total_number_reviews, Reviews, Bought With]
Index: []

duplicated skus: 3026

Summary For Stock Data - After Cleaning

	Count	Missing	Empty	Unique	Type	String	Int	Float	\
sku	1154	0	0	1154	object	1154	0	0	
stock_name	1154	0	0	1135	object	1154	0	0	
description	1154	0	0	1081	object	1154	0	0	
list_price	1154	0	0	322	float64	0	0	1154	
stock_image_url	1154	0	0	1141	object	1154	0	0	
brand	1154	0	0	228	object	1154	0	0	

	List
sku	0
stock_name	0
description	0
list_price	0
stock_image_url	0
brand	0

First 3 Rows

	sku	stock_name	\
0	0903a80	KitchenAid® Artisan® 5-qt. Stand Mixer KSM150PS	
1	13cab12	JCPenney Home Saratoga Cut-to-Width Fringed B...	
2	13e154b	Glamorise® Full-Figure Body Briefer - 6201	

	description	list_price	\
0	The mixer you've always dreamed of. Unique mix...	604.31	
1	Saratoga cut--to-width blackout shade features...	27.80	
2	Glamorise's best-selling full-figure body brie...	81.97	

	stock_image_url	brand
0	http://s7d9.scene7.com/is/image/JCPenney/09006...	Kitchen Aid
1	http://s7d2.scene7.com/is/image/JCPenney/DP121...	JCP HOME
2	http://s7d9.scene7.com/is/image/JCPenney/09006...	Glamorise

1.2.6 Sales

Details of all sales activity. Contains 7,982 sales records with each uniquely identified by the key 'uniq_id'. The data for each sale includes the sales price, stock reference and sales channel information. Most of the data appeared complete and reasonably, although several hundred sales prices were in an invalid format or missing; the relatively small number should not skew later analysis.

Data Content After review and validation the created dataframe, sales_df, has 7,982 records. It consists of:

- uniq_id - A unique identifier for the sales activity
- sku - A cross-reference for stock data in the stock_df
- sale_price - The price that the sales was
- category_tree - A string breaking down the structure of the sales channel
- category - The bottom level of the category tree
- sales_product_url - JCP website url for the product details as sold
- average_product_rating - An average of the customer review scores (1 to 5) for this sale
- total_number_reviews - The total number of customer reviews for this sale
- bought_with_list - other sales at the same time as this sale

Collation & Validation The provided data file, jcpenny_products.json, was examined and all sales specific data extracted into the sales_df dataframe. The following actions are highlighted:

- Invalid & Missing Prices: 263 sales prices were in a range format (34.5-45.9) and these were converted taking the average. And 18 had no price and so were zeroed *Lookup list price in stock?*
- Categories Missing: 636 categories, category trees are missing. About 10% of the 7,982 sales
- Sales Product URL: All good, no missing or duplicated *use request to test a sampe of URLS*
- Bought With: *check uniq-ids match and are in the main list what information does it provide*

```
[7]: # Establish a list of all product sales

# Use the initial sales file creating during the stock file creation,
# and drop columns that were retained in the stock_df
columns_not_required = ['name_title', 'description', 'list_price',
↳ 'product_image_urls', 'brand']
sales_df.drop(columns=columns_not_required, inplace=True, errors='ignore')

# Print the file and data fields
print_full_summary('Summary For Product Data - Initial Look', sales_df)

# Rename the retained columns (nb all listed for documentation purposes)
sales_df = sales_df.rename(columns={'uniq_id': 'uniq_id',
                                  'sku': 'sku',
                                  'sale_price': 'sale_price',
                                  'category': 'category', 'category_tree':
↳ 'category_tree',
                                  'average_product_rating':
↳ 'average_product_rating',
                                  'product_url': 'sales_product_url',
                                  'total_number_reviews':
↳ 'total_number_reviews',
```

```

        'Reviews': 'reviews_list', 'Bought With':
        ↪ 'bought_with_list'})

# Flag all missing fields for easier checking and replacement
missing_flag = 'Missing'
sales_df.replace('', missing_flag, inplace=True)
sales_df.fillna(missing_flag, inplace=True)

# uniq_id - nb visual inspection shows no duplicate and no missing
# sku - validated during production of the sales file

# sale_price - several values missing and formatted as range numbers
# Clean up and convert to float
def convert_price(price):
    try:
        # Trap the values with a range
        if '-' in price:
            low, high = map(float, price.split('-'))
            averaged = (low + high) / 2
            return averaged
        if price == missing_flag:
            return 0.0
        return float(price)
    except:
        return 0.0
sales_df['sale_price'] = sales_df['sale_price'].apply(convert_price)

# category and category tree
missing_cat = len(sales_df[sales_df['category'] == missing_flag])
missing_cat_tree = len(sales_df[sales_df['category_tree'] == missing_flag])
print(f'Missing: Categories {missing_cat} Trees {missing_cat_tree}')

# product URL check
duplicates_count = sales_df.duplicated(subset=['sales_product_url'],
        ↪ keep=False).sum()
print(f'Duplicated URLs: {duplicates_count}')

# Create an initial new dataframe for customer reviews ready for later
        ↪ manipulation
working_customer_reviews_df = sales_df.copy()
columns_not_required = ['sku', 'sale_price', 'category', 'category_tree',
        ↪ 'sales_product_url', 'bought_with_list']
working_customer_reviews_df.drop(columns=columns_not_required, inplace=True,
        ↪ errors='ignore')

# From this sales df, drop reviews details, but keep averages

```

```

# (nb the average and totals will be cross-checked as part of creating the
↳ reviews df)
columns_not_required = ['reviews_list']
sales_df.drop(columns=columns_not_required, inplace=True, errors='ignore')

'''
non_empty = group[group['stock_image_url'] != missing_flag]
if(len(non_empty) != 0):
    most_frequent = non_empty['stock_image_url'].value_counts().idxmax()
    new_sku.at[0, 'stock_image_url'] = most_frequent
else:
    new_sku.at[0, 'stock_image_url'] = 'No URL Available'
'''

# Print the file and data fields
print_full_summary('Summary For Sales Data - After Cleaning', sales_df)

# Tidy up
del columns_not_required
del missing_flag, missing_cat_tree, missing_cat, duplicates_count

```

Summary For Product Data - Initial Look

	Count	Missing	Empty	Unique	Type	String	Int	\
uniq_id	7982	0	0	7982	object	7982	0	
sku	7982	0	0	6110	object	7982	0	
sale_price	7982	0	0	2063	object	7982	0	
category	7982	0	0	1169	object	7982	0	
category_tree	7982	0	0	1997	object	7982	0	
average_product_rating	7982	0	0	153	float64	0	0	
product_url	7982	0	0	7982	object	7982	0	
total_number_reviews	7982	0	0	22	int64	0	7982	
Reviews	7982	0	0	7982	object	0	0	
Bought With	7982	0	0	7982	object	0	0	

	Float	List
uniq_id	0	0
sku	0	0
sale_price	0	0
category	0	0
category_tree	0	0
average_product_rating	7982	0
product_url	0	0
total_number_reviews	0	0
Reviews	0	7982
Bought With	0	7982

First 3 Rows

	uniq_id	sku	sale_price	category	\
0	b6c0b6bea69c722939585baeac73c13d	pp5006380337	24.16	alfred dunner	
1	93e5272c51d8cce02597e3ce67b7ad0a	pp5006380337	24.16	alfred dunner	
2	013e320f2f2ec0cf5b3ff5418d688528	pp5006380337	24.16	view all	

	category_tree	average_product_rating	\
0	jcpenny women alfred dunner	2.625	
1	jcpenny women alfred dunner	3.000	
2	jcpenny women view all	2.625	

	product_url	total_number_reviews	\
0	http://www.jcpenny.com/alfred-dunner-essentia...	8	
1	http://www.jcpenny.com/alfred-dunner-essentia...	8	
2	http://www.jcpenny.com/alfred-dunner-essentia...	8	

	Reviews	\
0	[{'User': 'fsdv4141', 'Review': 'You never hav...	
1	[{'User': 'tpcu2211', 'Review': 'You never hav...	
2	[{'User': 'pcfg3234', 'Review': 'You never hav...	

	Bought With
0	[898e42fe937a33e8ce5e900ca7a4d924, 8c02c262567...
1	[bc9ab3406dcaa84a123b9da862e6367d, 18eb69e8fc2...
2	[3ce70f519a9cfdd85cdbdec358e5347, b0295c96d2b...

Missing: Categories 636 Trees 636

Duplicated URLs: 0

Summary For Sales Data - After Cleaning

	Count	Missing	Empty	Unique	Type	String	Int	\
uniq_id	7982	0	0	7982	object	7982	0	
sku	7982	0	0	6110	object	7982	0	
sale_price	7982	0	0	1992	float64	0	0	
category	7982	0	0	1169	object	7982	0	
category_tree	7982	0	0	1997	object	7982	0	
average_product_rating	7982	0	0	153	float64	0	0	
sales_product_url	7982	0	0	7982	object	7982	0	
total_number_reviews	7982	0	0	22	int64	0	7982	
bought_with_list	7982	0	0	7982	object	0	0	

	Float	List
uniq_id	0	0
sku	0	0
sale_price	7982	0
category	0	0
category_tree	0	0
average_product_rating	7982	0
sales_product_url	0	0
total_number_reviews	0	0

bought_with_list 0 7982

First 3 Rows

	uniq_id	sku	sale_price	category \
0	b6c0b6bea69c722939585baeac73c13d	pp5006380337	24.16	alfred dunner
1	93e5272c51d8cce02597e3ce67b7ad0a	pp5006380337	24.16	alfred dunner
2	013e320f2f2ec0cf5b3ff5418d688528	pp5006380337	24.16	view all

	category_tree	average_product_rating \
0	jcpenny women alfred dunner	2.625
1	jcpenny women alfred dunner	3.000
2	jcpenny women view all	2.625

	sales_product_url	total_number_reviews \
0	http://www.jcpenny.com/alfred-dunner-essentia...	8
1	http://www.jcpenny.com/alfred-dunner-essentia...	8
2	http://www.jcpenny.com/alfred-dunner-essentia...	8

	bought_with_list
0	[898e42fe937a33e8ce5e900ca7a4d924, 8c02c262567...
1	[bc9ab3406dcaa84a123b9da862e6367d, 18eb69e8fc2...
2	[3ce70f519a9cfdd85cdbdec358e5347, b0295c96d2b...

1.2.7 Customer Sales Reviews

There are a total of 39,063 reviews but only 29,464 appear to be unique review comments. Further analysis found that 15,535 (40%) of reviews were used by several customers, worst case being several instances of 18 customers using the same comments. This could be because the sample data has been automatically generated or that customer ids are being created to generate false reviews. This data has *not* been dropped from the dataset, although later analysis of the reviews could be misleading.

Data Content After review and validation the created dataframe, `customer_reviews_df`, has 39,063 customer reviews. It consists of:

- `uniq_id` - A cross-reference for sales data in the `sales_df`
- `customer_id` - A cross-reference for customer data in the `customers_df`
- `review_text` - Review comments made by the customer for the sale
- `review_score` - The score of 1 to 5 given by the customer

Collation & Validation The provided data file, `jcpenny_products.json`, was examined and all review specific data extracted into the `customer_reviews_df` dataframe. This was carried out after first creating the `stock_df` and the `sales_df`. The following actions were taken:

- Reviews: From the sales details the list of customer reviews was decoded from its JSON format
- Ratings: The totals and averages across multiple customers were cross-checked to the sales data

- Customers: 17 reviews did not have a valid existing customer and a dummy customer was created (but flagged with state ISO of XX and DOB of NAT) so as not to lose the review data
- In 11 instances there were two reviews for the same customer and sale, but with different review comments; these were left as only a small number

CSV File Rejected The reviews.csv file was examined and the scores were found to have a large number of zero values (11,265 out of 39,063) and a quick examination showed that many scores differ between the JSON and CSV source. This confirmed the decision to reject the reviews.CSV data and only use the JSON source.

```
[8]: # Establish the dataframe with all customer reviews
# (nb initial working df prepared earlier during validation of the sales df)

customer_reviews_df = pd.DataFrame(columns=['uniq_id', 'customer_id',
    ↪ 'review_text', 'review_score'])

# Print the file and data fields
print_full_summary('Summary For Customer Reviews Data - Reviews Not Decoded',
    ↪ working_customer_reviews_df)

# Extract all reviews held in JSON format in the reviews list
# Include the uniq_id to link each review back to the original sales details

# Iterate through all rows of the original products data, to extract and decode
    ↪ the series of JSON items
# TODO: Itertuples is fast but this takes 4 seconds to run, replace with a
    ↪ more efficient approach?
for row in working_customer_reviews_df.itertuples(index=False):
    # Create reviews from decoded JSON & add the uniq_id for this row
    reviews_json = row.reviews_list
    reviews_dict_string = json.dumps(reviews_json)
    reviews_df = pd.DataFrame(json.loads(reviews_dict_string))
    reviews_df.columns = ['customer_id', 'review_text', 'review_score']
    reviews_df.insert(0, 'uniq_id', row.uniq_id)
    reviews_df = reviews_df

    # Cross-check the customer_id for each review
    # Several not found so need to iterate through each and create a dummy
    ↪ customer record
    if not reviews_df.customer_id.isin(customers_df.customer_id).all():
        print(f'For: {row.uniq_id} not all customers match')
        for customer in reviews_df.itertuples(index=False):
            if customer.customer_id not in customers_df.customer_id.values:
                print(f'Adding dummy customer for ID: {customer.customer_id}')
                new_customer = pd.DataFrame([{'customer_id': customer.
    ↪ customer_id,
```

```

        'DOB': pd.NaT,
        'state_ISO': 'XX',
        'uniq_id_list': []]])
    customers_df = pd.concat([customers_df, new_customer],
    ignore_index=True)

    # Check the average and totals originally in the source file
    if not math.isclose(row.average_product_rating, reviews_df.review_score.
    mean()):
        print (f'For: {row.uniq_id} ratings mismatch, original: {row.
    average_product_rating} vs {reviews_df.review_score.mean()}')
    if not math.isclose(row.total_number_reviews, len(reviews_df)):
        print (f'For: {row.uniq_id} counts mismatch, original: {row.
    total_number_reviews} vs {len(reviews_df)}')

    # Add the review to the customers review df
    customer_reviews_df = pd.concat([customer_reviews_df, reviews_df])

    # Print the completed file and data fields
    print_full_summary('Summary For Customer Reviews Data - All Reviews',
    customer_reviews_df)

    # Tidy up
    del row, reviews_json, reviews_dict_string, reviews_df,
    working_customer_reviews_df, new_customer, customer

```

Summary For Customer Reviews Data - Reviews Not Decoded

	Count	Missing	Empty	Unique	Type	String	Int	\
uniq_id	7982	0	0	7982	object	7982	0	
average_product_rating	7982	0	0	153	float64	0	0	
total_number_reviews	7982	0	0	22	int64	0	7982	
reviews_list	7982	0	0	7982	object	0	0	

	Float	List
uniq_id	0	0
average_product_rating	7982	0
total_number_reviews	0	0
reviews_list	0	7982

First 3 Rows

	uniq_id	average_product_rating	\
0	b6c0b6bea69c722939585baeac73c13d	2.625	
1	93e5272c51d8cce02597e3ce67b7ad0a	3.000	
2	013e320f2f2ec0cf5b3ff5418d688528	2.625	

	total_number_reviews	reviews_list
0	8	[{'User': 'fsdv4141', 'Review': 'You never hav...

```

1          8  [{'User': 'tpcu2211', 'Review': 'You never hav...
2          8  [{'User': 'pcfg3234', 'Review': 'You never hav...

```

For: e5bdf53f2374569526c9f4d55afdd88e not all customers match

Adding dummy customer for ID: dqft3311

Summary For Customer Reviews Data - All Reviews

	Count	Missing	Empty	Unique	Type	String	Int	Float	\
uniq_id	39063	0	0	7982	object	39063	0	0	
customer_id	39063	0	0	4993	object	39063	0	0	
review_text	39063	0	0	29464	object	39063	0	0	
review_score	39063	0	0	5	object	0	39063	0	

	List
uniq_id	0
customer_id	0
review_text	0
review_score	0

First 3 Rows

	uniq_id	customer_id	\
0	b6c0b6bea69c722939585baeac73c13d	fsdv4141	
1	b6c0b6bea69c722939585baeac73c13d	krpz1113	
2	b6c0b6bea69c722939585baeac73c13d	mbmg3241	

	review_text	review_score
0	You never have to worry about the fit...Alfred...	2
1	Good quality fabric. Perfect fit. Washed very ...	4
2	I do not normally wear pants or capris that ha...	4

```

[9]: # Further cross-validation of reviews and customers data

# Tidy up customers_df to cross-check and drop the uniq_id_list

# Check that all sales listed in the customers df match customer reviews
# !! This takes 24 seconds, so very slow, but needs to iterate nested to handle
↳ individual exceptions
for customer in customers_df.itertuples(index=False):
    sales_list = customer.uniq_id_list
    for uniq_id in sales_list:
        # Check if the sale exists in the sales df
        if uniq_id in sales_df.uniq_id.values:
            # Check that a unique review exists in the customer reviews df for
↳ the sale + customer
            matched_reviews = customer_reviews_df.loc[
                (customer_reviews_df['customer_id'] == customer.customer_id) &
                (customer_reviews_df['uniq_id'] == uniq_id)]
            if len(matched_reviews) != 1:

```

```

        print(f'For customer: {customer.customer_id} + uniq_id_{
↳{uniq_id} reviews = {len(matched_reviews)}')
    else:
        print(f'For customer: {customer.customer_id}. A uniq_id {uniq_id}_{
↳does not exist}')

# Drop the uniq_id_list from the customers df as all covered in the customer_{
↳reviews file
customers_df.drop(columns='uniq_id_list', inplace=True, errors='ignore')

# Tidy up
del customer, sales_list, uniq_id, matched_reviews

```

```

For customer: fwbl1442 + uniq_id fe4541f4c1dde497edda95fa46e9e98d reviews = 2
For customer: DUP001dqft3311 + uniq_id 5f280fb338485cfc30678998a42f0a55 reviews
= 0
For customer: ffxf2322 + uniq_id b28c5fe83b8b20b05c2451e79cea85f1 reviews = 2
For customer: vwuj3242 + uniq_id e7bea081cac88a6bdcbb1d447a4253bab reviews = 2
For customer: ntvh2341 + uniq_id fedc1fca14619493cd14436a9817c4f2 reviews = 2
For customer: slos2412 + uniq_id 63251a30df90f586fb769ddf2aa5ed54 reviews = 2
For customer: mbdt1413 + uniq_id 77661aaf8abd87167e310721616c6f6a reviews = 2
For customer: pawj4231 + uniq_id a60d13f2f6313bd961546c40c6a3ca96 reviews = 2
For customer: jeph4124 + uniq_id 2dcd61eaea3a7ded2049f305391ae2b8 reviews = 2
For customer: DUP002dqft3311 + uniq_id 571b86d307f94e9e8d7919b551c6bb52 reviews
= 0
For customer: ndkl1344 + uniq_id 387d1795d7221b01252a2d8eff30ba87 reviews = 2
For customer: fnmd4431 + uniq_id 07647adc11b605d1a50ccc163eb96c54 reviews = 2
For customer: wnm2211 + uniq_id 6f7a799e8e5bd4c959379217a776eb86 reviews = 2

```

```

[10]: # Look at how many reviews are duplicates and how many customers are linked to_{
↳these

duplicates_by_customer = customer_reviews_df.
↳groupby('review_text')['customer_id'].size().reset_index(name='cust_count')
reviews_duplicated = duplicates_by_customer.groupby('cust_count').count().
↳reset_index()

count_reviews_single = reviews_duplicated[reviews_duplicated['cust_count'] == 1]
count_reviews_duplicated = len(customer_reviews_df) -_{
↳count_reviews_single['review_text'].sum()
max_duplicates = reviews_duplicated['cust_count'].max()

print(f'Out of a total of {len(customer_reviews_df)} reviews_{
↳{count_reviews_duplicated} are duplicates.')
print(f'Or approximately {(count_reviews_duplicated/len(customer_reviews_df))_{
↳* 100):.0f}%')

```

```

print(f'Several worst case situations with {max_duplicates} customers using the_
↳same review comments.')

# Tidy up
del duplicates_by_customer
del reviews_duplicated
del count_reviews_single
del count_reviews_duplicated
del max_duplicates

```

Out of a total of 39063 reviews 15535 are duplicates.

Or approximately 40%

Several worst case situations with 18 customers using the same review comments.

```

[11]: # Load the CSV reviews file to cross-check against the data extracted from the_
↳JSON sourced reviews

# Load the reviews .csv file, exit if do not exist or are invalid
file_path = os.path.join(os.getcwd(), DATA_DIRECTORY, 'reviews.csv')
if not os.path.isfile(file_path):
    raise Exception(f"File not found: {file_path}")
source_reviewsCSV_df = pd.read_csv(file_path)

# Initial look at the file and data fields
print(f'Summary for customer reviews - CSV')
print_file_summary(source_reviewsCSV_df)
print(f'First 3 rows')
display(source_reviewsCSV_df.head(3))

# Scores look very different
count_zero_scores = source_reviewsCSV_df[source_reviewsCSV_df['Score'] ==_
↳0]['Score'].count()
count_zero_scoresJSON = customer_reviews_df[customer_reviews_df['review_score']_
↳== 0]['review_score'].count()

print(f'Compare JSON sourced review vs CSV file source')
print(f'Count: {len(customer_reviews_df)} vs {len(source_reviewsCSV_df)}')
print(f'Scores with zero: {count_zero_scoresJSON:.0f} vs {count_zero_scores:_
↳0f}')
print(f'Mean: {customer_reviews_df['review_score'].mean():.1f} vs_
↳{source_reviewsCSV_df['Score'].mean():.1f}')

# Tidy Up
del count_zero_scores, count_zero_scoresJSON, source_reviewsCSV_df

```

Summary for customer reviews - CSV

Count	Missing	Empty	Unique	Type	String	Int	Float	List
-------	---------	-------	--------	------	--------	-----	-------	------

Uniq_id	39063	0	0	7982	object	39063	0	0	0
Username	39063	0	0	4993	object	39063	0	0	0
Score	39063	0	0	6	int64	0	39063	0	0
Review	39063	0	0	29463	object	39063	0	0	0

First 3 rows

	Uniq_id	Username	Score	\
0	b6c0b6bea69c722939585baeac73c13d	fsdv4141	2	
1	b6c0b6bea69c722939585baeac73c13d	krpz1113	1	
2	b6c0b6bea69c722939585baeac73c13d	mbmg3241	2	

Review

0	You never have to worry about the fit...Alfred...
1	Good quality fabric. Perfect fit. Washed very ...
2	I do not normally wear pants or capris that ha...

Compare JSON sourced review vs CSV file source

Count: 39063 vs 39063

Scores with zero: 0 vs 11265

Mean: 3.0 vs 1.5

1.2.8 Tidy Up and Save

```
[12]: # Store the completed working dataframes for analysis in seperate workbook
%store sales_df stock_df customer_reviews_df customers_df states_df
```

```
Stored 'sales_df' (DataFrame)
Stored 'stock_df' (DataFrame)
Stored 'customer_reviews_df' (DataFrame)
Stored 'customers_df' (DataFrame)
Stored 'states_df' (DataFrame)
```